E-PANNs: An Efficient version of PANNs for Audio Tagging

ARSHDEEP SINGH[1], HAOHE LIU, MARK D. PLUMBLEY

CENTRE FOR VISION, SPEECH AND SIGNAL PROCESSING (CVSSP), UNIVERSITY OF SURREY, UK.

[1]RESEARCH FELLOW A, EMAIL: ARSHDEEP.SINGH@SURREY.AC.UK.

SPEECH AND AUDIO IN THE NORTHEAST (SANE) WORKSHOP 2023, NEW YORK UNIVERSITY (NYU), BROOKLYN, NEW YORK.

INTRODUCTION 1

- Pre-trained audio neural networks (PANNs) [1] are a family of convolutional neural networks (CNNs) designed for audio classification tasks.
- PANNs have shown state-of-the-art performance in audio tagging and have been widely employed for many downstream tasks including audio scene classification and several DCASE [2] related tasks as a feature extractor or end-to-end classifier.
- However, the best performing PANNs require 21G computations* (multiply-accumulate operations, MACs) for inference and requires 81M parameters for storage.
- Large size and computations demand more resources and slow down the inference speed.
- This work contribute to reduce computations and memory storage requirement of PANNs while maintaining the performance as given by the original PANNs.

FILTER PRUNING TO OBTAIN E-PANNS 3

- **Filter Pruning [3]** involves "removing" some of the filters and their connected feature maps/channels from convolutional neural networks (CNNs) to compress them (See Figure 2 (a)).
- **Benefits of Pruning:** Reduces computational cost (multiply-accumulate operations, MACs) and memory storage of CNNs while providing similar performance, hence CNNs become more efficient.

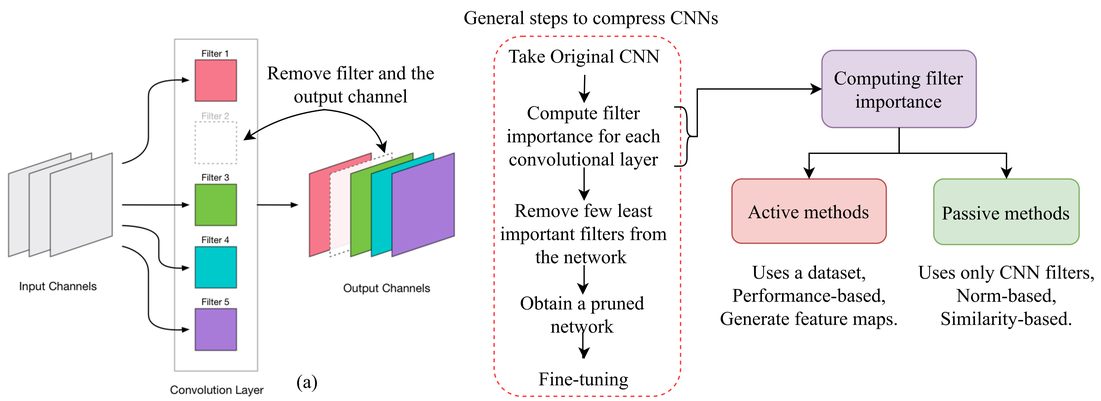

Figure 2

We opt a passive filter pruning approach to compute filter importance.A filter producing significant output, as measured using operator norm of the filter [4,5], is considered more important than others.

REFERENCES

[1] Qiuqiang Kong et al., "PANNs: Large-scale pretrained audio neural networks for audio pattern recognition", IEEE/ACM Transactions on Audio, Speech, and Language Processing, 28:2880–2894, 2020.

[2] Detection and Classification of Acoustic Scenes and Events (DCASE) challenge (https://dcase.community/)

[3] Luo, Jian-Hao, et al. "ThiNet: pruning CNN filters for a thinner net", IEEE transactions on pattern analysis and machine intelligence 41.10 (2018): 2525-2538.

[4] Arshdeep Singh , H Liu, and Mark D. Plumbley (2023) "E-PANNs: Sound Recognition Using Efficient Pre-trained Audio Neural Networks", 52nd International Congress and Exposition on Noise Control Engineering (Inter-Noise 2023), Chiba, Greater Tokyo, Japan, 20-23 August 2023.

[5] Singh, A, and Mark D. Plumbley. "Efficient CNNs via Passive Filter Pruning", arXiv preprint arXiv:2304.02319 (2023).

[6] Our Live Sound recognition demo available at: https://github.com/Arshdeep-Singh-Boparai/E-PANNs

OUR CONTRIBUTION 2

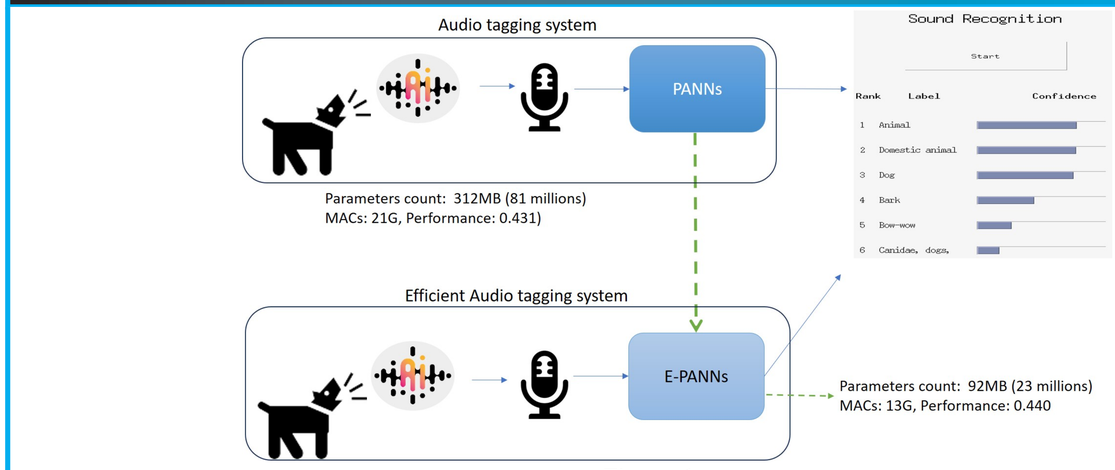

Figure 1

PERFORMANCE ANALYSIS & SUMMARY 4

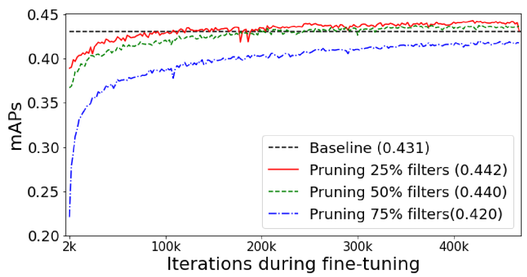

Figure 3

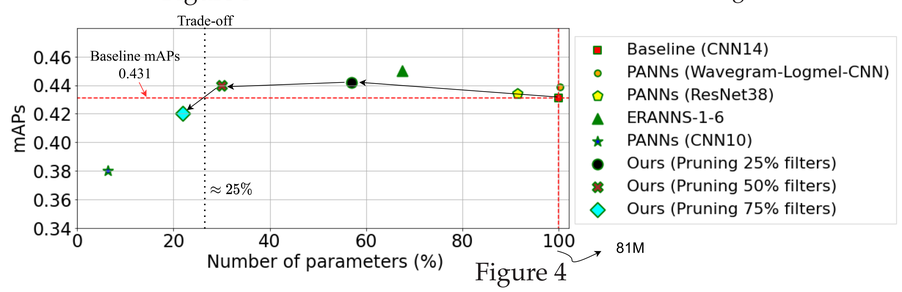

Figure 4

Summary

- A pre-trained network is made efficient by removing unimportant filters followed by a fine-tuning process (with 20-50% efforts as used in training).
- Reduced 30% computations, 70% memory storage with an improved performance for audio tagging (sound recognition) task.
- Live sound recognition demo is available at [6].

ACKNOWLEDGMENTS

This work was partly supported by Engineering and Physical Sciences Research Council (EPSRC) Grant EP/T019751/1 "AI for Sound (AI4S)". For the purpose of open access, the authors have applied a Creative Commons Attribution (CC BY) licence to any Author Accepted Manuscript version arising.

Personal profile: https://www.surrey.ac.uk/people/arshdeep-singh